

Use of GPU resources for the Training of Neural Networks

Tim Voigtländer, Roger Wolf, Lars Sowa
tim.voigtlaender@student.kit.edu

Karlsruher Institut für Technologie (KIT) - Institut für experimentelle Teilchenphysik (ETP)



Will a GPU speed up your Machine Learning application?

It's complicated

- GPU speed up training immensely by parallelizing demanding tasks
 - Some overhead is necessary to transfer data
 - Individual computations are also a bit slower
 - Usually a GPU replaces multiple CPU cores
- Hard to determine if an ML application will be sped up

Rule of thumb: GPUs help most for

- Bigger networks
- Less complex networks

There are many hyperparameters to tune


The only sure way to know, is to try it out



Example: Event Discrimination via NN


- To differentiate between signal and background, Neural Networks are trained
 - 22 to 28 physical variables are used as inputs
 - 5 classes are used as outputs
 - 4 are composed of various backgrounds
 - The signal for a specific group of mass hypotheses
 - The network is comparably simple
 - Feed forward fully connected network
 - 3 hidden layers with 512 nodes each
- Around **45,000** parameters

InceptionResnetV2 has around **56,000,000** Parameters




Example: CPU vs GPU

- **Part of Master thesis:**
 - Adapt training for GPU
 - Find ideal Hyperparameters
 - Compare with CPU
- **CPU only**
 - 12 CPU cores
 - Average runtime: 29 minutes
- **CPU + GPU**
 - 2 CPU cores + A100 GPU
 - Average runtime: 7.5 Minutes



➤ Inclusion of GPU leads to a significant speedup while using fewer CPU resources

What so we need for CPU training




How a training on local CPU might look like (**executable .sh**)

```
#!/bin/bash


# Source all necessary software
source /work/name/some/local/environment/setup.sh

# Run the training
python3 /work/name/some/local/files/mnist_training.py
--files /ceph/name/some/local/files/training_data
```

What so we need for GPU training



What so we need for GPU training



What GPUs can you use?

The “Deepthought” machine

- Local development machine
- Ubuntu 18.04
- 4 GTX TITAN X

The TOpAS cluster

- Remotely usable through the ETP batch system
- CentOS7
- Multiple machines with
 - V100, V100S, A100
- 56 GPUs in total

What GPUs can you use?

The “Deepthought” machine

- Local development machine
- Ubuntu 18.04
- 4 GTX TITAN X

The TOpAS cluster

- Remotely usable through the ETP batch system
- CentOS7
- Multiple machines with
 - V100, V100S, A100
- 56 GPUs in total

What is in the Jobfile for a CPU job

```
# The Executable that is run in the Job  
  
Executable = ./executable.sh  
...  
  
# The hardware requirements  
  
RequestCPUs = 1          # 1 CPU core  
  
RequestMemory = 2000      # 2 GB of memory  
  
request_disk = 2000000    # 2 GB of disk space  
...  
...
```

[HTCondor Talk by Matthias Schnepf](#)

[Corresponding GitLab repo](#)

What is in the Jobfile for a GPU job on TOpAS

...

```
## Additional Arguments in submission.jdl

# Request GPU resources

# Can also be >1 to request more GPUs

request_GPUs = 1

# Mark as remote job

+RemoteJob = True
```

What is in the Jobfile for a GPU job on TOpAS

...

```
## Additional Arguments in submission.jdl

# Request GPU resources

# Can also be >1 to request more GPUs

request_GPUs = 1

# Mark as remote job

+RemoteJob = True
```

What is a Remote Job?

Remote jobs do not have access to `/ceph`, `/work` and `/home`

What is a Remote Job?

Remote jobs do not have access to **/ceph**, **/work** and **/home**

Neither

`source /work/name/some/local/environment/setup.sh`


nor

`python3 /work/name/some/local/files/mnist_training.py`

`--files /ceph/name/some/local/files/training_data`

will work!

What so we need for GPU training on TOpAS



How to transfer Files to Remote Batch Jobs

Small files (Together < 1 GB):

- Through HTCondor internal transfer

[Shown in GitLab examples](#)

How to transfer small files

```
...  
  
## Additional Arguments in submission.jdl  
  
# List of files to transfer into job  
  
transfer_input_files = mnist_training.py  
  
# List of files to transfer out of job  
  
transfer_output_files = trained_model.pt  
  
# Paths can be adjusted with transfer_output_remaps
```

How to transfer Files to Remote Batch Jobs

Small files (Together < 1 GB):

- Through HTCondor internal transfer

[Shown in GitLab examples](#)

Big Files:

- Transfer with Grid protocols (SRM, XRootD, etc.)
 - Directly from `/ceph` via `/ceph/srv`
 - Can be slow
 - Through intermediate Grid storage
 - Often requires VOMS proxy

[Shown in GitLab examples](#)

How to transfer big files

1. Place files on accessible storage before the job starts

- a. `/ceph/srv` works for ETP and TOpAS
- b. Remote storage like NRG

How to transfer big files

1. Place files on accessible storage before the job starts
 - a. `/ceph/srv` works for ETP and TOpAS
 - b. Remote storage like NRG
2. Copy files into the job with Grid Protocols at runtime (like `xrdcp`)

Files on

`/ceph/srv/name/path/to/dataset`

can be found as

`root://ceph-node-a.etp.kit.edu://name/path/to/dataset`

[Documentation of XRootD](#)

How to transfer Files to Remote Batch Jobs

Small files (Together < 1 GB):

- Through HTCondor internal transfer


[Shown in GitLab examples](#)

Big Files:

- Transfer with Grid protocols (SRM, XRootD, etc.)
 - Directly from `/ceph` via `/ceph/srv`
 - Can be slow
 - Through intermediate Grid storage
 - Often requires VOMS proxy
- Some data types can be streamed
- Subject to caching on TOpAS

[Shown in GitLab examples](#)

What so we need for GPU training on TOpAS



Requirements for your GPU Environment

Set-up has to be:

- GPU capable

CUDA



cuDNN



**GPU versions
of software**

- Portable

Send with job

**Get from
CVMFS**

Get from Web

GPU Environment Examples: LCG Stack

```
source /cvmfs/.../LCG_101cuda/.../setup.sh
HOME=$_CONDOR_JOB_IWD
python3 -m venv venv
source venv/bin/activate
python3 -m pip install pip --upgrade
python3 -m pip install torch==1.9.0+cu111
    torchvision==0.10.0+cu111 torchaudio==0.9.0
    -f https://download.pytorch.org/whl/torch\_stable.html
```

[List of available LCG stacks](#)

Source base environment from CUDA LCG stacks

GPU Environment Examples: LCG Stack

```
source /cvmfs/.../LCG_101cuda/.../setup.sh
HOME=$_CONDOR_JOB_IWD
python3 -m venv venv
source venv/bin/activate
python3 -m pip install pip --upgrade
python3 -m pip install torch==1.9.0+cu111
    torchvision==0.10.0+cu111 torchaudio==0.9.0
    -f https://download.pytorch.org/whl/torch\_stable.html
```

[List of available LCG stacks](#)

Set `$HOME` to writable directory. `$_CONDOR_JOB_IWD` is the job start directory

GPU Environment Examples: LCG Stack

```
source /cvmfs/.../LCG_101cuda/.../setup.sh
HOME=$_CONDOR_JOB_IWD
python3 -m venv venv
source venv/bin/activate
python3 -m pip install pip --upgrade
python3 -m pip install torch==1.9.0+cu111
torchvision==0.10.0+cu111 torchaudio==0.9.0
-f https://download.pytorch.org/whl/torch\_stable.html
```

[List of available LCG stacks](#)

Set up local **venv** to install missing modules

GPU Environment Examples: LCG Stack

```
source /cvmfs/.../LCG_101cuda/.../setup.sh
HOME=$_CONDOR_JOB_IWD
python3 -m venv venv
source venv/bin/activate
python3 -m pip install pip --upgrade
python3 -m pip install torch==1.9.0+cu111
torchvision==0.10.0+cu111 torchaudio==0.9.0
-f https://download.pytorch.org/wheel/torch_stable.html
```

[List of available LCG stacks](#)

Install modules compatible with GPU Hardware and LCG stack libraries

GPU Environment Examples: LCG Stack

```
source /cvmfs/.../LCG_101cuda/.../setup.sh
```

```
HOME=$_CONDOR_JOB_IWD
```

```
python3 -m venv venv
```

```
source venv/bin/activate
```

```
python3 -m p...
```

Newest available CUDA version is 11.2

```
python3 -m p...
```

```
torchvision==0.10.0+cu111 torchaudio==0.9.0
```

```
-f https://download.pytorch.org/whl/torch\_stable.html
```

[List of available LCG stacks](#)

Install modules compatible with GPU Hardware and LCG stack libraries

GPU Environment Examples: Anaconda

```
HOME=$_CONDOR_JOB_IWD
wget https://repo.anaconda.com/miniconda/
    Miniconda3-latest-Linux-x86_64.sh
chmod +x Miniconda3-latest-Linux-x86_64.sh
./Miniconda3-latest-Linux-x86_64.sh -b -p $(pwd)/miniconda
source miniconda/bin/activate
conda install pytorch torchvision torchaudio
    cudatoolkit=11.3 -c pytorch -y
```

Set `$HOME` to writable directory. `$_CONDOR_JOB_IWD` is the job start directory

GPU Environment Examples: Anaconda

```
HOME=$_CONDOR_JOB_IWD  
wget https://repo.anaconda.com/miniconda/  
Miniconda3-latest-Linux-x86_64.sh  
chmod +x Miniconda3-latest-Linux-x86_64.sh  
. ./Miniconda3-latest-Linux-x86_64.sh -b -p $(pwd)/miniconda  
source miniconda/bin/activate  
conda install pytorch torchvision torchaudio  
cudatoolkit=11.3 -c pytorch -y
```

Download an Anaconda installer

GPU Environment Examples: Anaconda

```
HOME=$_CONDOR_JOB_IWD
wget https://repo.anaconda.com/miniconda/
    Miniconda3-latest-Linux-x86_64.sh
chmod +x Miniconda3-latest-Linux-x86_64.sh
./Miniconda3-latest-Linux-x86_64.sh -b -p $(pwd)/miniconda
source miniconda/bin/activate
conda install pytorch torchvision torchaudio
    cudatoolkit=11.3 -c pytorch -y
```

Set up Anaconda environment

GPU Environment Examples: Anaconda

```
HOME=$_CONDOR_JOB_IWD
wget https://repo.anaconda.com/miniconda/
Miniconda3-latest-Linux-x86_64.sh
chmod +x Miniconda3-latest-Linux-x86_64.sh
./Miniconda3-latest-Linux-x86_64.sh -b -p $(pwd)/miniconda
source miniconda/bin/activate
conda install pytorch torchvision torchaudio
cudatoolkit=11.3 -c pytorch -y
```

Install software compatible with GPU Hardware

GPU Environment Examples: Anaconda

```
HOME=$_CONDOR_JOB_IWD
wget https://repo.anaconda.com/miniconda/
    Miniconda3-latest-Linux-x86_64.sh
chmod +x Miniconda3-latest-Linux-x86_64.sh
./Miniconda3-latest-Linux-x86_64.sh
source miniconda3/bin/activate
miniconda
Newest available CUDA version is 11.3
conda install pytorch torchvision torchaudio
    cudatoolkit=11.3 -c pytorch -y
```

Install software compatible with GPU Hardware

GPU Environment Examples: Docker

Batch jobs at ETP are usually set up to run in a docker container

...

```
Universe = docker

## Docker image used by the job
# Default container

docker_image = mschnepf/slc7-condocker
# Container with CUDA + cuDNN libraries preinstalled

docker_image =
    twoigtlaender/slc7-condocker-cuda-11.6-cudnn8-devel:base
```

GPU Environment Examples: Docker

Batch jobs at ETP are usually set up to run in a docker container

...

```
Universe = docker

## Docker image used by the job

# Default container

docker_image = mschnepf/slc7-condocker

# Container with CUDA + cuDNN libraries preinstalled

docker_image =

twoigtlaender/slc7-condocker-cuda-11.6-cudnn8-devel:base
```

GPU Environment Examples: Docker

```
# CUDA and cuDNN are already installed from the start
HOME=$_CONDOR_JOB_IWD
python3 -m venv venv
source venv/bin/activate
python3 -m pip install pip --upgrade
python3 -m pip install torch==1.9.0+cu111
    torchvision==0.10.0+cu111 torchaudio==0.9.0
    -f https://download.pytorch.org/whl/torch\_stable.html
```

Set `$HOME` to writable directory. `$_CONDOR_JOB_IWD` is the job start directory

GPU Environment Examples: Docker

```
# CUDA and cuDNN are already installed from the start
HOME=$_CONDOR_JOB_IWD
python3 -m venv venv
source venv/bin/activate
python3 -m pip install pip --upgrade
python3 -m pip install torch==1.9.0+cu111
torchvision==0.10.0+cu111 torchaudio==0.9.0
-f https://download.pytorch.org/whl/torch\_stable.html
```

Set up local **venv** to install missing modules

GPU Environment Examples: Docker

```
# CUDA and cuDNN are already installed from the start
HOME=$_CONDOR_JOB_IWD
python3 -m venv venv
source venv/bin/activate
python3 -m pip install pip --upgrade
python3 -m pip install torch==1.9.0+cu111
    torchvision==0.10.0+cu111 torchaudio==0.9.0
    -f https://download.pytorch.org/whl/torch_stable.html
```

Install modules compatible with GPU Hardware and preinstalled libraries

GPU Environment Examples: Docker

```
# CUDA and cuDNN are already installed from the start
HOME=$_CONDOR_JOB_IWD
python3 -m venv venv
source venv/bin/activate
python3 -m pip install torch torchvision torchaudio
python3 -m pip install -f https://download.pytorch.org/whl/torch_stable.html
torchvision==0.10.0+cu111 torchaudio==0.9.0
```

Newest available CUDA version is 11.6

```
-f https://download.pytorch.org/whl/torch_stable.html
```

Install modules compatible with GPU Hardware and preinstalled libraries

GPU Environment Examples: `/cvmfs/etp.kit.edu`

If you have a stable environment setup, it can be put on

`/cvmfs/etp.kit.edu`

Files there can be **read** from inside the batch job

[CVMFS Documentation](#)

GPU Environment Examples: `/cvmfs/etp.kit.edu`

If you have a stable environment setup, it can be put on

`/cvmfs/etp.kit.edu`

Files there can be **read** from inside the batch job


[CVMFS Documentation](#)

The **Anaconda** environment setup placed on `/cvmfs/etp.kit.edu`:

`source`

`/cvmfs/etp.kit.edu/GPU_examples/miniconda/bin/activate`
GPU_ML

What so we need for GPU training on TOpAS



How to train on GPU: Some Notes

- The example is a shortened version of the GitLab example ([Link](#))
- Everything was tested on fairly new versions (PyTorch **1.9**, Tensorflow **2.4**)
 - Significantly older versions (especially Tensorflow 1) will deviate
 - GPU training with Tensorflow 1 is still possible but requires more work
- Ideally, analyses should not run on old software
- I can provide examples for older versions if necessary

The Training on CPU (`mnist_training.py`)

```
# Define NN model

class Net(torch.nn.Module) :

    < Definition of model >

model = Net()

# Other setups

optimizer, scheduler,
train_loader, validation_loader
```

The Training on CPU ([mnist_training.py](#))

```
# Define NN model

class Net(torch.nn.Module) :

    < Definition of model >

model = Net()

# Other setups

optimizer, scheduler,
train_loader, validation_loader
```

The Training on CPU (`mnist_training.py`)

```
# Run training loop
for epoch in range(1, epochs + 1):
    train(model, ...)
    validation(model, ...)
    scheduler.step()

# Export trained model
torch.save(model.state_dict(), out_dir)
```

The Training on CPU (`mnist_training.py`)

```
# Run training loop
for epoch in range(1, epochs + 1):
    train(model, ...)
    validation(model, ...)
    scheduler.step()

# Export trained model
torch.save(model.state_dict(), out_dir)
```

The Training on CPU (`mnist_training.py`)

```
# The training step

def train(model, train_loader, optimizer):
    model.train()

    for data, target in train_loader:

        optimizer.zero_grad()
        output = model(data)
        loss = LossFunctional(output, target)
        loss.backward()
        optimizer.step()
```

The Training on CPU (`mnist_training.py`)

```
# The validation step

def validation(model, validation_loader):

    model.eval()

    loss = 0

    with torch.no_grad():

        for data, target in validation_loader:

            output = model(data)

            loss += LossFunctional(output, target)
```

How to train on GPU with PyTorch

1. Detect the GPU

2. Place model on GPU

3. Place data on GPU

[Article on GPU usage in PyTorch](#)

How to train on GPU with PyTorch

1. Detect the GPU
2. Place model on GPU
3. Place data on GPU

[Article on GPU
usage in PyTorch](#)

How to train on GPU with PyTorch

1. Detect the GPU
2. Place model on GPU
3. Place data on GPU

[Article on GPU usage in PyTorch](#)

How to train on GPU with PyTorch

1. Detect the GPU
 2. Place model on GPU
 3. Place data on GPU
- PyTorch commands utilize the GPU

[Article on GPU usage in PyTorch](#)

How to detect the GPU

```
# Detect if GPU are usable
GPU_check = torch.cuda.is_available()

# Detect how many GPUs are usable
num_GPUs = torch.cuda.device_count()

# Set which device to use
device = torch.device("cuda" if GPU_check
else "cpu")
```

How to detect the GPU

```
# Detect if GPU are usable
GPU_check = torch.cuda.is_available()

# Detect how many GPUs are usable
num_GPUs = torch.cuda.device_count()

# Set which device to use
device = torch.device("cuda" if GPU_check
else "cpu")
```

How to detect the GPU

```
# Detect if GPU are usable
GPU_check = torch.cuda.is_available()

# Detect how many GPUs are usable
num_GPUs = torch.cuda.device_count()

# Set which device to use
device = torch.device("cuda" if GPU_check
else "cpu")
```

The Training on CPU (`mnist_training.py`)

```
# Define NN model

class Net(torch.nn.Module) :

    < Definition of model >

model = Net()

# Other setups

optimizer, scheduler,
train_loader, validation_loader
```

The Training on GPU (`mnist_training.py`)

```
# Define NN model

class Net(torch.nn.Module) :

    < Definition of model >

model = Net()

model = model.to(device)
```

- On CPU
- On Device

```
# Other setups

optimizer, scheduler,
train_loader, validation_loader
```

The Training on CPU (`mnist_training.py`)

```
# Run training loop
for epoch in range(1, epochs + 1):
    train(model, ...)
    validation(model, ...)
    scheduler.step()

# Export trained model
torch.save(model.state_dict(), out_dir)
```

The Training on CPU (`mnist_training.py`)

```
# Run training loop
for epoch in range(1, epochs + 1):
    train(model, device, ...)
    validation(model, device, ...)
    scheduler.step()
```

- On CPU
- On Device

```
# Export trained model
torch.save(model.state_dict(), out_dir)
```

The Training on CPU (`mnist_training.py`)

```
# The training step

def train(model, train_loader, optimizer):
    model.train()

    for data, target in train_loader:

        optimizer.zero_grad()
        output = model(data)
        loss = LossFunctional(output, target)
        loss.backward()
        optimizer.step()
```

The Training on GPU (`mnist_training.py`)

```
# The training step

def train(model, device, train_loader, optimizer):
    model.train()

    for data, target in train_loader:
        data, target = data.to(device), target.to(device)

        optimizer.zero_grad()

        output = model(data)

        loss = LossFunctional(output, target)

        loss.backward()

        optimizer.step()
```

- On CPU
- On Device

The Training on CPU (`mnist_training.py`)

```
# The validation step

def validation(model, validation_loader):

    model.eval()

    loss = 0

    with torch.no_grad():

        for data, target in validation_loader:

            output = model(data)

            loss += LossFunctional(output, target)
```

The Training on CPU (`mnist_training.py`)

```
# The validation step

def validation(model, device, validation_loader):

    model.eval()
    loss = 0
    with torch.no_grad():

        for data, target in validation_loader:

            data, target = data.to(device), target.to(device)

            output = model(data)

            loss += LossFunctional(output, target)
```

- On CPU
- On Device

How to train on GPU with TensorFlow

- TensorFlow commands utilize the GPU by default

[Guide to GPU usage
with TensorFlow](#)

How to train on GPU with TensorFlow

- TensorFlow commands utilize the GPU by default
- GPUs can still be managed manually

[Guide to GPU usage with TensorFlow](#)

```
# Detect which GPUs are available
```

```
GPU_list = tensorflow.config.list_physical_devices('GPU')
```

```
# Run code on specific device
```

```
with tensorflow.device(GPU_list[0]):  
    < Some calculation >
```

How to train on GPU with TensorFlow

- TensorFlow commands utilize the GPU by default
- GPUs can still be managed manually

[Guide to GPU usage with TensorFlow](#)

```
# Detect which GPUs are available
```

```
GPU_list = tensorflow.config.list_physical_devices('GPU')
```

```
# Run code on specific device
```

```
with tensorflow.device(GPU_list[0]):  
    < Some calculation >
```

GPU training on TOpAS: Putting it all together

Submit remote GPU job via `condor_submit submit.jdl`

GPU training on TOpAS: Putting it all together

Submit remote GPU job via `condor_submit submit.jdl`

What should happen in the Executable of the job:

1. Set up GPU capable environment

GPU training on TOpAS: Putting it all together

Submit remote GPU job via `condor_submit submit.jdl`

What should happen in the Executable of the job:

1. Set up GPU capable environment
2. Copy training data into job

GPU training on TOpAS: Putting it all together

Submit remote GPU job via `condor_submit submit.jdl`

What should happen in the Executable of the job:

1. Set up GPU capable environment
2. Copy training data into job
3. Run GPU adapted NN training

GPU training on TOpAS: Putting it all together

Submit remote GPU job via `condor_submit submit.jdl`

What should happen in the Executable of the job:

1. Set up GPU capable environment
2. Copy training data into job
3. Run GPU adapted NN training
4. Return trained model from job

Related Topics

- Partial- and Multi GPU usage
- (GPU)-Job management
- GPU related Hyperparameter tuning
- Framework specific performance optimization
- More about Hardware related improvements for ML
- ML Profiling with Tensorboard

**Let us know if you are interested
in any of these topics**

Collection of Links

- [GitLab repo with examples for this talk](#)
- [Indico of Matthias Schnepfs HTCondor talk](#)
 - [Corresponding GitLab repo](#)
- [Documentation of XRootD](#)
- [List of available LCG stacks](#)
- [CVMFS Documentation](#)
- [Article on GPU usage in PyTorch](#)
- [Guide to GPU usage with TensorFlow](#)
- [TOpAS monitoring](#)

Thank you for
your attention