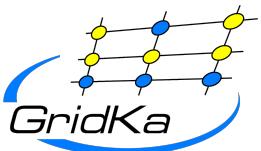




# HTCondor Tutorial

Matthias Schnepf | 17. July 2025



# Computing Resources at ETP

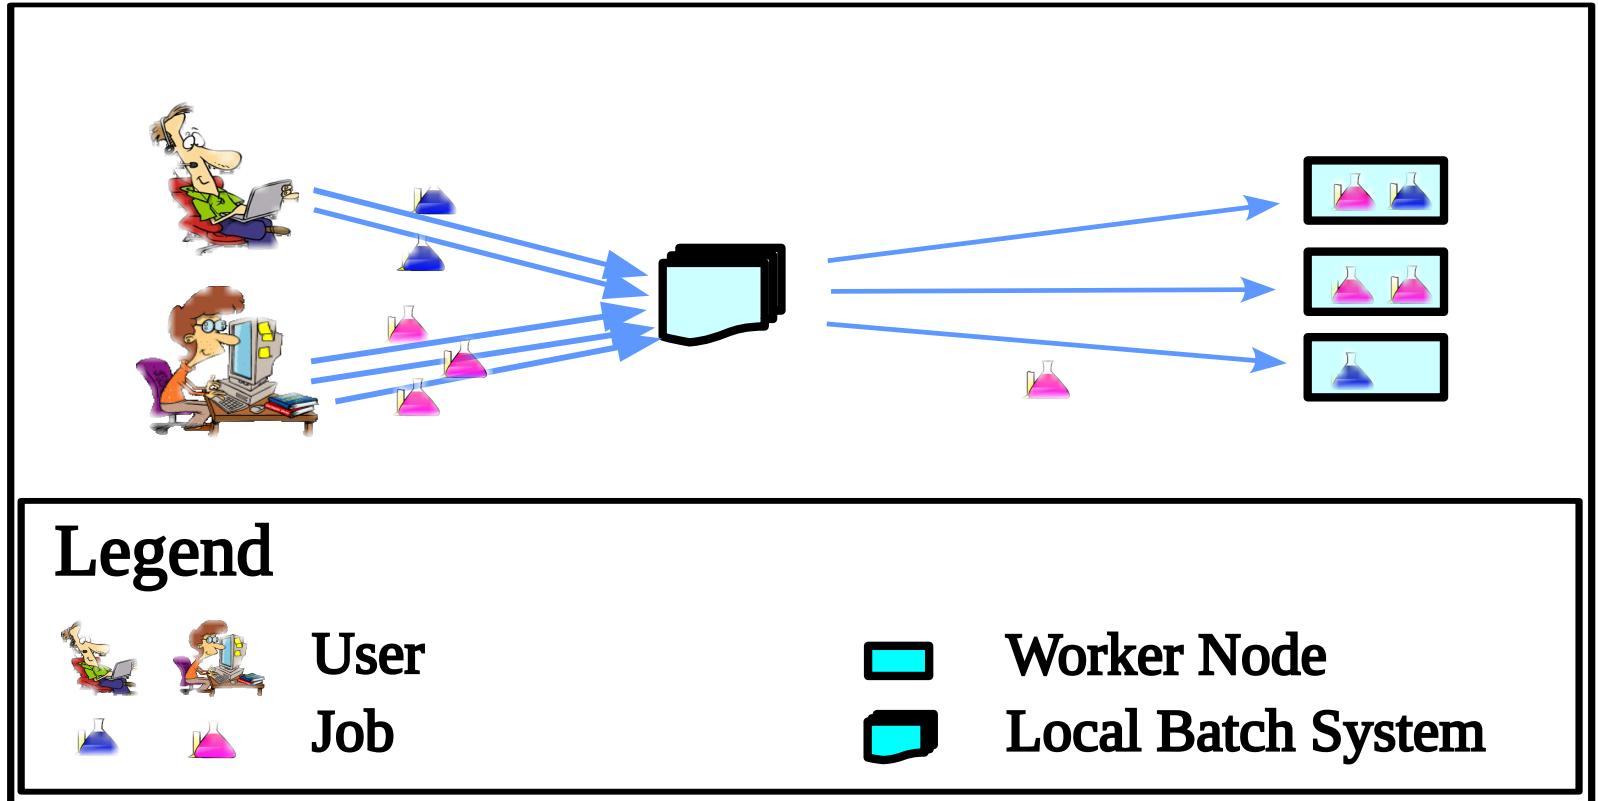
- end-user device: laptop/desktop PC/tablet,...
  - web browsing
  - video conference
  - login to other computing resources
- portal machines: bms[1,3] portal1, deepthought(GPUs)
  - powerful machines to develop and test software
  - can be used for **special** workflows
  - access to batch system/Grid
- batch system/Grid
  - distributed bunch of machines (worker nodes)
  - majority of computing resources

# Storage Resources

- storage mounted on all desktops, portal machines and some of the worker nodes:
  - */home*
    - for configs, mails, ...
    - 50 GB limit per user
    - backup every second day
  - */work*
    - for software and files with fast access
    - 200 GB limit per user
    - weekly backup
  - */ceph*
    - for large datasets
    - about 1.3 PB total
    - please clean up from time to time
    - no backup
- Grid storage
  - for datasets and software sandboxes
  - accessible from everywhere via Grid protocols (see later)
  - collaboration storage
    - official dataset
  - NRG (**National Grid Resource Germany**)
    - for local group
    - located at GridKa with  $80 \text{ Gbit s}^{-1}$  network connection to ETP
    - CMS: 2.75 PB
    - Belle II 250 TB

# Batch System

- batch system: coordinated run of batch jobs (non-interactive programs) on distributed resources
  - submit nodes (portal machines)
  - worker nodes
  - central node
- batch job
  - executable
  - meta data
    - name of executable
    - memory requirements
    - submit host
    - worker node
    - ...



Based on "The Pilot Way to Grid Resources Using glideinWMS" from I. Sfiligoi et. al.

# Batch Job (1)

- executable (program/script that run on a worker node)

1. (worker node information)
2. setup environment
3. copy data to worker node if necessary
4. run workload
5. copy output to storage e.g. /ceph or NRG

```
#!/bin/bash

##### WN information
echo -e "Hostname:$(hostname)"
echo -e "user:$(whoami)"
echo -e "spawndir:$(pwd)"

##### setup environment
source /cvmfs/belle.cern.ch/tools/b2setup release-09-00-01

##### start workload at spawndir and create output.root with results
basf2 /work/$(whoami)/HTCondor_Tutorial/code/01_simple_submit/mc_example.py -o output.root

##### copy output from WN
cp output.root /work/$(whoami)/HTCondor_Tutorial/code/01_simple_submit/simple_submit_output.root
```

# Batch Job (1)

- executable (program/script that run on a worker node)

1. (worker node information)
2. setup environment
3. copy data to worker node if necessary
4. run workload
5. copy output to storage e.g. /ceph or NRG

- should be **tested** on portal machine before submission

- check for errors
- check for resource demand via htop: CPU, number of processes, Memory,...

```
#!/bin/bash

##### WN informaton
echo -e "Hostname:$(hostname)"
echo -e "user:$(whoami)"
echo -e "spawndir:$(pwd)"

##### setup environment
source /cvmfs/belle.cern.ch/tools/b2setup release-09-00-01

##### start workload at spawndir and create output.root with results
basf2 /work/$(whoami)/HTCondor_Tutorial/code/01_simple_submit/mc_example.py -o output.root

##### copy output from WN
cp output.root /work/$(whoami)/HTCondor_Tutorial/code/01_simple_submit/simple_submit_output.root
```

# Batch Job (1)

- executable (program/script that run on a worker node)

1. (worker node information)
2. setup environment
3. copy data to worker node if necessary
4. run workload
5. copy output to storage e.g. /ceph or NRG

- should be **tested** on portal machine before submission

- check for errors
- check for resource demand via htop: CPU, number of processes, Memory,...

- develop portable code

- job runtime should be longer than 20 min

- output file size should be bigger than 1 GB

```
#!/bin/bash

##### WN information
echo -e "Hostname:$(hostname)"
echo -e "user:$(whoami)"
echo -e "spawndir:$(pwd)"

##### setup environment
source /cvmfs/belle.cern.ch/tools/b2setup release-09-00-01

##### start workload at spawndir and create output.root with results
basf2 /work/$(whoami)/HTCondor_Tutorial/code/01_simple_submit/mc_example.py -o output.root

##### copy output from WN
cp output.root /work/$(whoami)/HTCondor_Tutorial/code/01_simple_submit/simple_submit_output.root
```

# Batch Job (2)

## ■ submit file (meta data about the job)

- container environment
- executable
- resource demand
  - CPU cores (default 1)
  - RAM (default 2000 MB)
  - Disk (default 500.000 kB)
  - Walltime (default 86 400 sec)
- logging
- accounting group
  - beamdump (5%)
  - belle (25%)
  - cms (70%)
- submit (queue)

## ■ submit with *condor\_submit submit\_file.jdl*

```
Universe = container
container_image = /cvmfs/unpacked.cern.ch/registry.hub.docker.com/cverstege/alma9-gridjob:latest
executable = ./simple_submit.sh
request_cpus = 1
request_memory = 2000
request_disk = 500000
+RequestWalltime = 86400
log = log.log
stdout = stdout.log
stderr = stderr.log
accounting_group=belle
queue
```

# Job Status

user

idle

hold

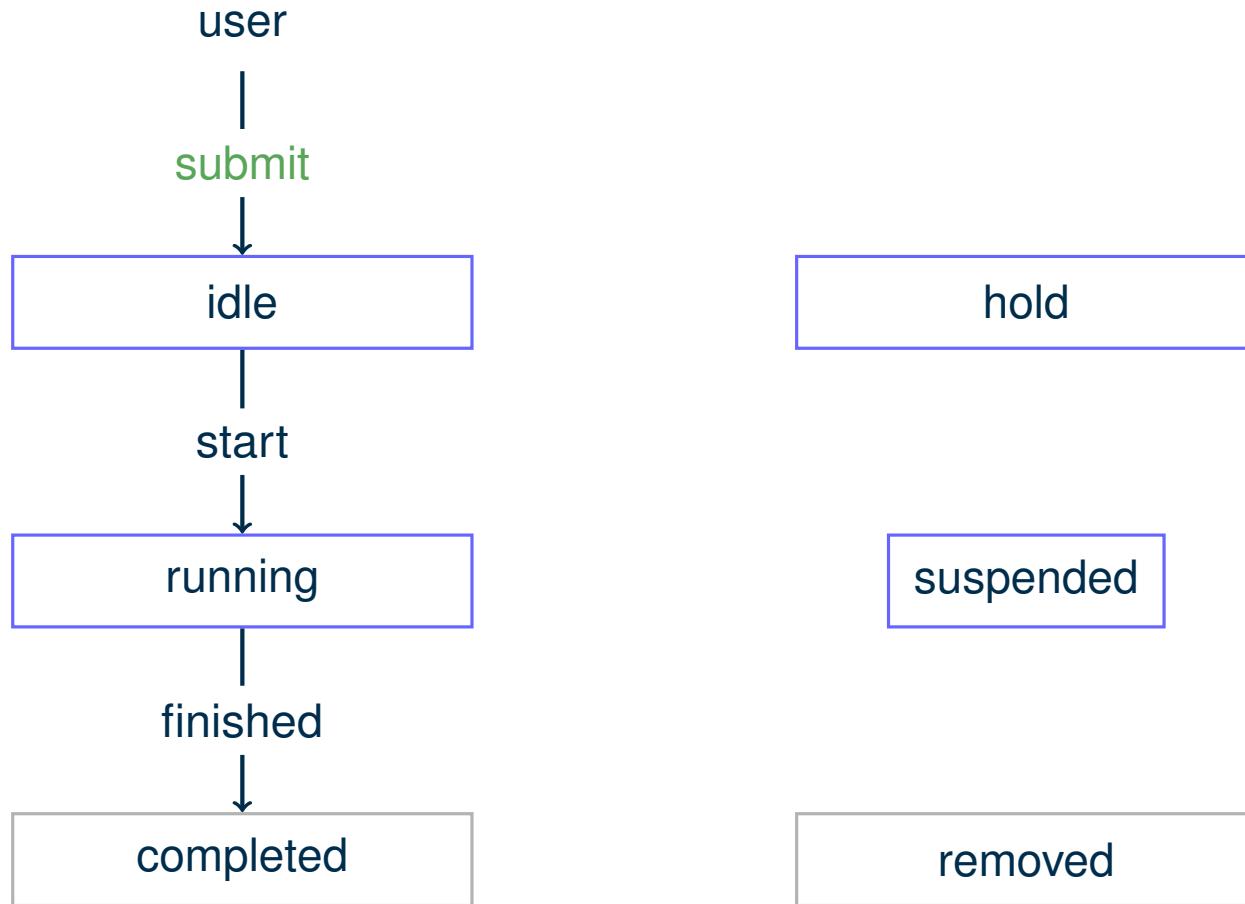
running

suspended

completed

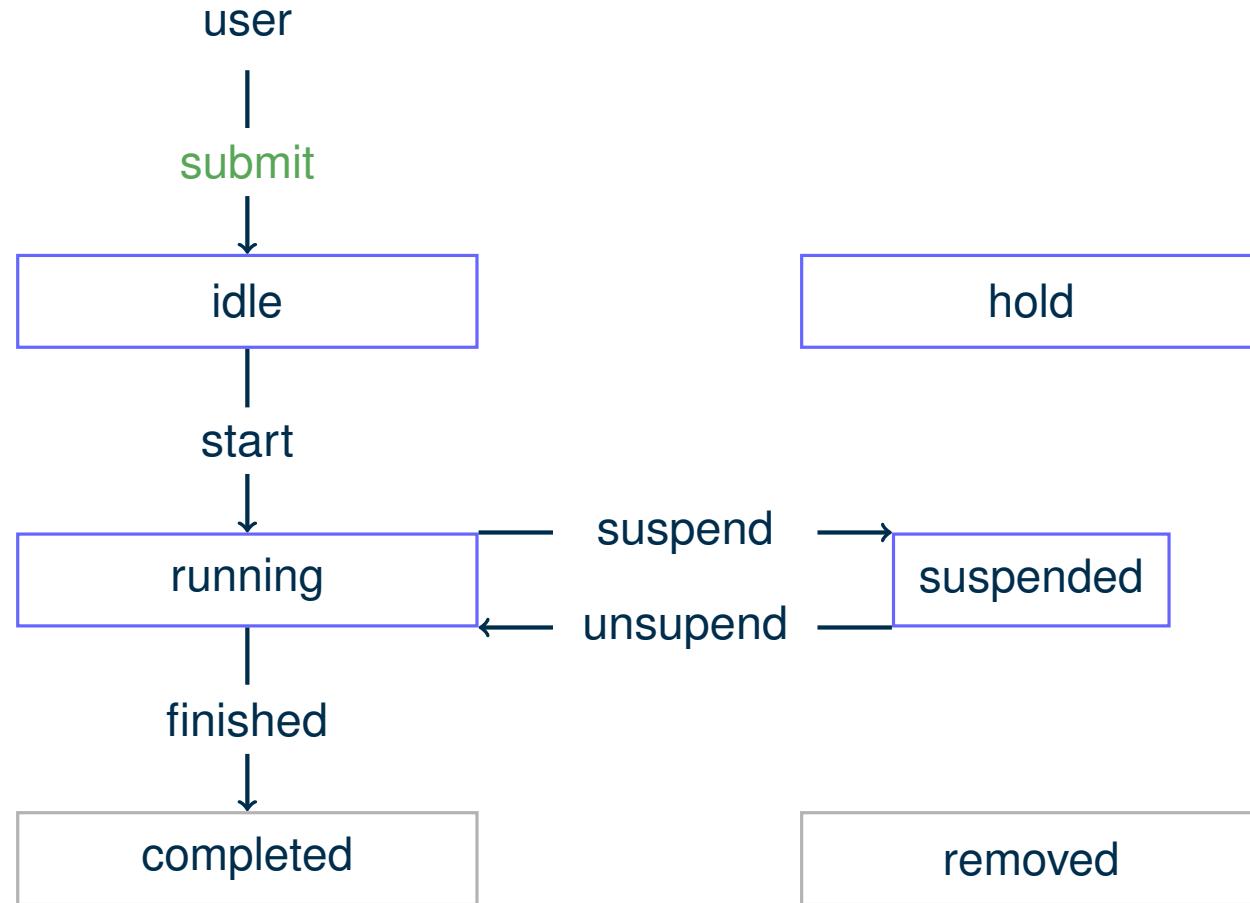
removed

# Job Status



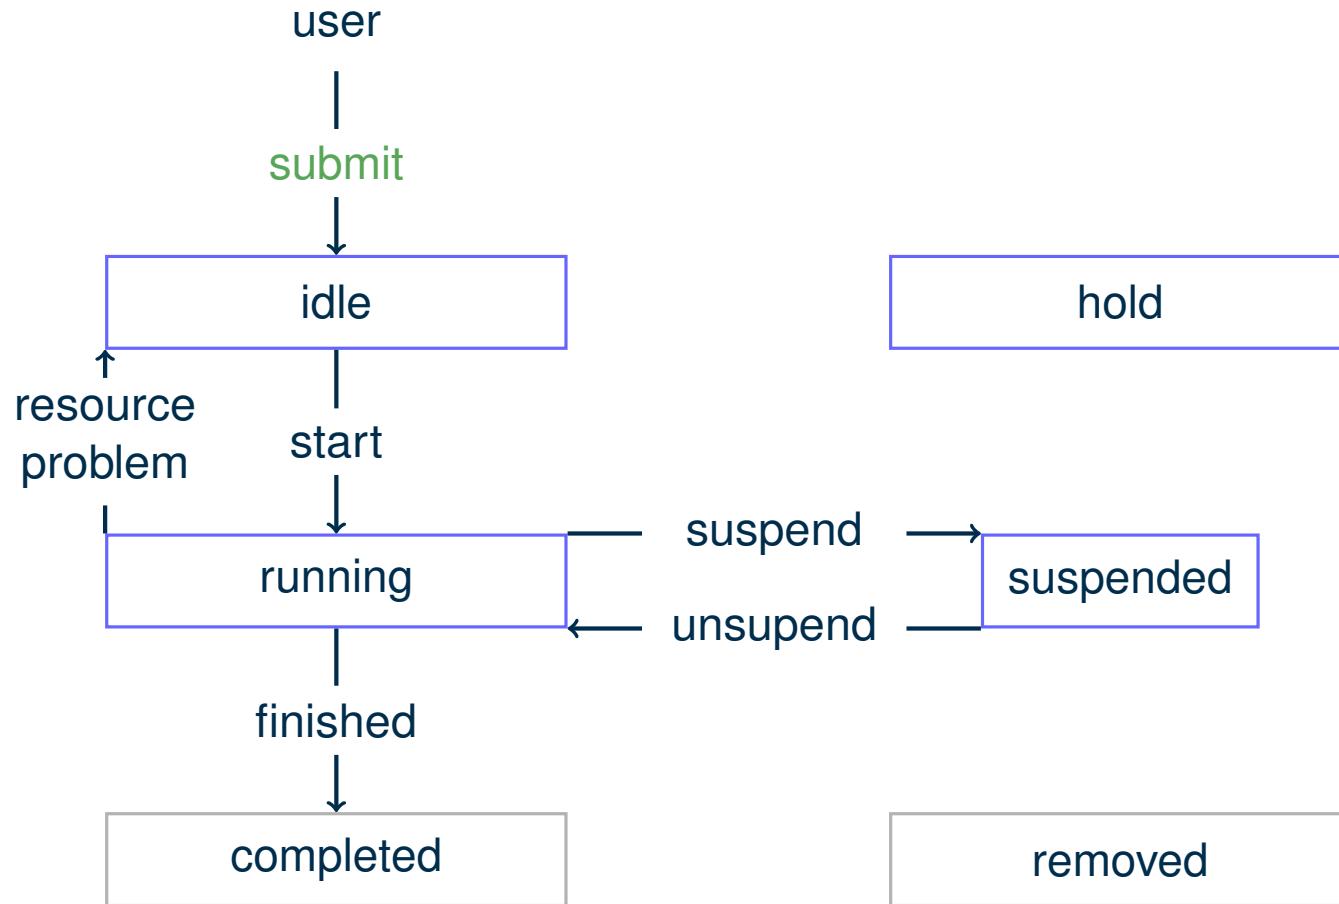
- only green action can be done by users

# Job Status



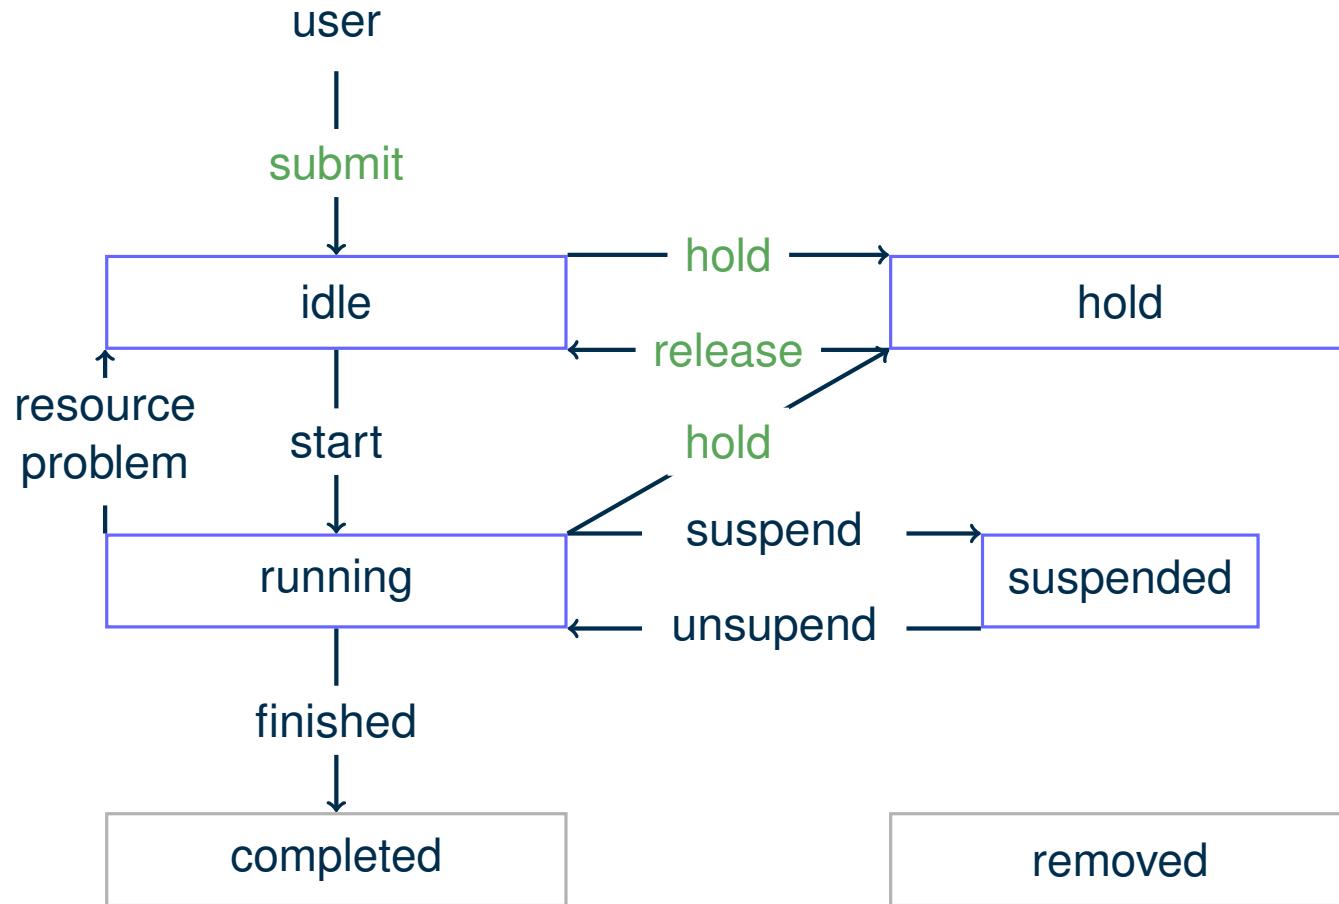
- only green action can be done by users
- only running and suspended jobs are on WNs

# Job Status



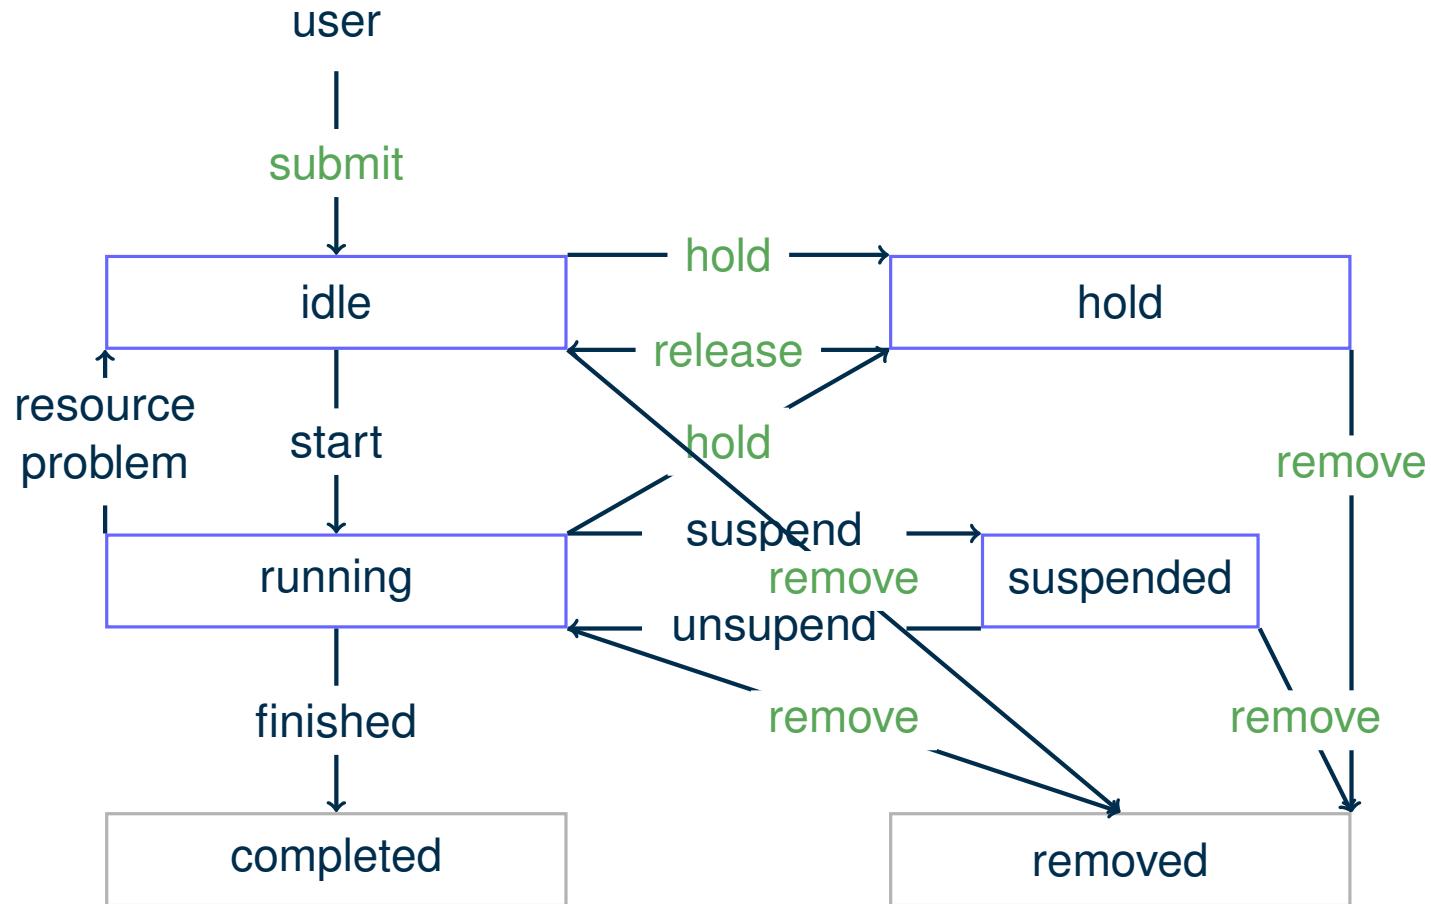
- only green action can be done by users
- only running and suspended jobs are on WNs

# Job Status



- only green action can be done by users
- only running and suspended jobs are on WNs
- hold jobs are removed from the WN, unsaved results are gone

# Job Status



- only green action can be done by users
- only running and suspended jobs are on WNs
- hold jobs are removed from the WN, unsaved results are gone
- completed and removed jobs are history and can not be changed

# Useful Condor Commands

- *condor\_submit*
  - submit job to queue of the batch system
- *condor\_q*
  - shows the status of jobs in queue
  - see jobs from all submit nodes use `-g <username>` option
- *condor\_rm*
  - remove jobs from queue
  - running jobs get killed and removed
- *condor\_history*
  - shows job history (removed and completed jobs)
  - ETP history is **long**; please use `-limit <number>`

# Useful Condor Commands

- *condor\_submit*
  - submit job to queue of the batch system
- *condor\_q*
  - shows the status of jobs in queue
  - see jobs from all submit nodes use `-g <username>` option
- *condor\_rm*
  - remove jobs from queue
  - running jobs get killed and removed
- *condor\_history*
  - shows job history (removed and completed jobs)
  - ETP history is **long**; please use `-limit <number>`
- *condor\_release*
  - change jobs status from *hold* to *idle*
- *condor\_q*, *condor\_rm*, *condor\_history*, *condor\_q\_edit* and *condor\_release* has similar selections
  - `<user>`
  - `ClusterID.ProcID / ClusterID`
  - `-constraint 'classadd expression'` e.g. `-constraint 'RemoteJob =?= True'`

# Useful Condor Commands

- *condor\_submit*
  - submit job to queue of the batch system
- *condor\_q*
  - shows the status of jobs in queue
  - see jobs from all submit nodes use `-g <username>` option
- *condor\_rm*
  - remove jobs from queue
  - running jobs get killed and removed
- *condor\_history*
  - shows job history (removed and completed jobs)
  - ETP history is **long**; please use `-limit <number>`
- *condor\_release*
  - change jobs status from *hold* to *idle*
- *condor\_q*, *condor\_rm*, *condor\_history*, *condor\_q\_edit* and *condor\_release* has similar selections
  - `<user>`
  - `ClusterID.ProcID / ClusterID`
  - `-constraint 'classadd expression'` e.g. `-constraint 'RemoteJob =?= True'`
- *condor\_q\_edit*
  - edit job attributes

# Useful Condor Commands

- *condor\_submit*
  - submit job to queue of the batch system
- *condor\_q*
  - shows the status of jobs in queue
  - see jobs from all submit nodes use `-g <username>` option
- *condor\_rm*
  - remove jobs from queue
  - running jobs get killed and removed
- *condor\_history*
  - shows job history (removed and completed jobs)
  - ETP history is `long`; please use `-limit <number>`
- *condor\_release*
  - change jobs status from *hold* to *idle*
- *condor\_q*, *condor\_rm*, *condor\_history*, *condor\_q\_edit* and *condor\_release* has similar selections
  - `<user>`
  - `ClusterID.ProcID / ClusterID`
  - `-constraint 'classadd expression'` e.g. `-constraint 'RemoteJob =?= True'`
- *condor\_q\_edit*
  - edit job attributes
- *condor\_status*
  - get information of the machines in the system
  - status of the WNs per default

# Useful Condor Commands

- *condor\_submit*
  - submit job to queue of the batch system
- *condor\_q*
  - shows the status of jobs in queue
  - see jobs from all submit nodes use `-g <username>` option
- *condor\_rm*
  - remove jobs from queue
  - running jobs get killed and removed
- *condor\_history*
  - shows job history (removed and completed jobs)
  - ETP history is `long`; please use `-limit <number>`
- *condor\_release*
  - change jobs status from *hold* to *idle*
- *condor\_q*, *condor\_rm*, *condor\_history*, *condor\_q\_edit* and *condor\_release* has similar selections
  - `<user>`
  - `ClusterID.ProcID / ClusterID`
  - `-constraint 'classadd expression'` e.g. `-constraint 'RemoteJob =?= True'`
- *condor\_q\_edit*
  - edit job attributes
- *condor\_status*
  - get information of the machines in the system
  - status of the WNs per default
- *condor\_userprio*
  - shows user prio per accounting group
  - `-allusers` shows prio of all users

# Job Information

- HTCondor works with information in form of ClassAds
- ClassAds are **key value/expression pairs** e.g.
  - `Owner = "mschnepf"`
  - `START = TARGET.Owner=?="mschnepf"`
  - `requirements = (TARGET.Memory >= RequestMemory)`
- every instance in HTCondor has ClassAds: Jobs, WNs, scheduler,...
  - get ClassAds of `jobs` via `condor_q/condor_history`
  - get ClassAds of `daemons` via `condor_status`
  - `-l` shows all ClassAds of the instance
  - `-af ClassAd` shows only the ClassAd value/expression

## ■ important ClassAds

- `requirements`: requirements of the job to slot
- `CloudSite`: group of worker nodes with similar setup
- `HoldReason`: reason why your job is on hold
- `RemoteUserCPU`: amount of CPU time the job used on a worker node

## ■ ClassAd syntax

- keys are not case sensitive, only strings in quotes are case sensitive
- values can be *undefined*
- is equal (`=?=`) returns only false or true while `==` can return false, true, undefined
- is not equal (`=!=`)
- and (`&&`), or (`||`)

# Submit a hand full of Jobs

- several jobs with the same executable but different arguments, e.g., input file or seed
- queue command enables to submit a cluster of jobs
  - each job has the same *ClusterID*
  - each job has another *ProcID*
  - Job ID is *ClusterID*.*ProcID*
  - *ClusterID* and *ProcID* can be used in submit file

```
Universe = container
container_image = /cvmfs/unpacked.cern.ch/registry.hub.docker.com/cverstege/alma9-gridjob:latest

executable = ./executable.sh
arguments = $(ProcID)

request_cpus = 1
request_memory = 2000
request_disk = 500000

log = log_$(ProcID).log
stdout = stdout_$(ProcID).log
stderr = stderr_$(ProcID).log

accounting_group=belle

queue 5
```

# Submit a hand full of Jobs

- several jobs with the same executable but different arguments, e.g., input file or seed
- queue command enables to submit a cluster of jobs
  - each job has the same *ClusterID*
  - each job has another *ProcID*
  - Job ID is *ClusterID*.*ProcID*
  - *ClusterID* and *ProcID* can be used in submit file
  - can read variables from file; one line per job

```
Universe = container
container_image = /cvmfs/unpacked.cern.ch/registry.hub.docker.com/cverstege/alma9-gridjob:latest
executable = ./executable.sh

request_cpus = 1
request_memory = 2000
request_disk = 500000

log = log_$(ProcID).log
stdout = stdout_$(ProcID).log
stderr = stderr_$(ProcID).log

accounting_group=belle

queue arguments from input.csv
```

# Interactive Job

- request an interactive job

- get a job slot for interactive use (no other program is running) with resources set in the submit file
- add *-interactive* to your submit command
- ends when you disconnect, e.g., via *exit*

- jump into your job with *condor\_ssh\_to\_job JOBID*

- see your running processes
- access to files of your job
- able to copy files from your job
- when jobs end, you get kicked out



generated with Microsoft Copilot

# Hold State

- you set the job to hold
- your job tried something and failed
  - used more memory/disk than requested
  - restarted more than 5 time
  - log directory does not exist
  - output file of job was not created and can not copy to submit host
- hold jobs are removed from the WN, unsafed data are lost



Stop! Polizei! von iStock

# Hold State

- you set the job to hold
- your job tried something and failed
  - used more memory/disk than requested
  - restarted more than 5 time
  - log directory does not exist
  - output file of job was not created and can not copy to submit host
- hold jobs are removed from the WN, unsafed data are lost
- fix your job
  - get hold reason: `condor_q JOBID -hold`
  - adjust ClassAds: `condor_qedit JOBID CLASSAD NEW_VALUE`
  - release job: `condor_release JOBID`



Stop! Polizei! von iStock

# Hold State

- you set the job to hold
- your job tried something and failed
  - used more memory/disk than requested
  - restarted more than 5 time
  - log directory does not exist
  - output file of job was not created and can not copy to submit host
- hold jobs are removed from the WN, unsafed data are lost
- fix your job
  - get hold reason: `condor_q JOBID -hold`
  - adjust ClassAds: `condor_qedit JOBID CLASSAD NEW_VALUE`
  - release job: `condor_release JOBID`
- You can release a job up to 5 times. After that you have to adjust `JobRunCount` via `condor_qedit`



Stop! Polizei! von iStock

# Memory Usage

- keep your memory usage low

# Memory Usage

- keep your memory usage low
- job is on hold (job did not finish)

# Memory Usage

- keep your memory usage low
- job is on hold (job did not finish)
- get reason: `condor_q JOBID -hold`

# Memory Usage

- keep your memory usage low
- job is on hold (job did not finish)
- get reason: `condor_q JOBID -hold`
  - HoldReason: *Error from ... : Docker job has gone over memory limit of 2048 Mb*



bradhouse.dev - How to handle Out of Memory (OOM) in C - 2023

# Memory Usage

- keep your memory usage low
- job is on hold (job did not finish)
- get reason: `condor_q JOBID -hold`
  - HoldReason: *Error from ... : Docker job has gone over memory limit of 2048 Mb*
  - adjust memory request: `condor_qedit JOBID RequestMemory 3000`
  - release job: `condor_release JOBID`



bradhouse.dev - How to handle Out of Memory (OOM) in C - 2023

# Memory Usage

- keep your memory usage low
- job is on hold (job did not finish)
- get reason: `condor_q JOBID -hold`
  - HoldReason: *Error from ... : Docker job has gone over memory limit of 2048 Mb*
  - adjust memory request: `condor_qedit JOBID RequestMemory 3000`
  - release job: `condor_release JOBID`
- automated adjustment in JDL file
  - memory increase by each start: `request_memory = 2000+(1000*NumJobStarts)`
  - automated release after run into memory limit: `periodic_release = (HoldReasonCode == 34)`



bradhouse.dev - How to handle Out of Memory (OOM) in C - 2023

# Memory Usage

- keep your memory usage low
- job is on hold (job did not finish)
- get reason: `condor_q JOBID -hold`
  - HoldReason: *Error from ... : Docker job has gone over memory limit of 2048 Mb*
  - adjust memory request: `condor_qedit JOBID RequestMemory 3000`
  - release job: `condor_release JOBID`
- automated adjustment in JDL file
  - memory increase by each start: `request_memory = 2000+(1000*NumJobStarts)`
  - automated release after run into memory limit: `periodic_release = (HoldReasonCode == 34)`
- more memory ⇒ less resources



bradhouse.dev - How to handle Out of Memory (OOM) in C - 2023

# Why does my job not run?

- batch systems has to be full ;-)
- usually jobs has to wait a few minutes before start
- monitoring website

# Why does my job not run?

- batch systems has to be full ;-)
- usually jobs has to wait a few minutes before start
- monitoring website
- *condor\_q -better-analyse JOBID* provides information about the job-resource matchmaking

# Why does my job not run?

- batch systems has to be full ;-)
- usually jobs has to wait a few minutes before start
- monitoring website
- *condor\_q -better-analyse JOBID* provides information about the job-resource matchmaking
- see available resources via *condor\_status -constraint 'partitionableSlot' -af:t Name CPUs Memory Disk GPUs*

# Why does my job not run?

- batch systems has to be full ;-)
- usually jobs has to wait a few minutes before start
- monitoring website
- *condor\_q -better-analyse JOBID* provides information about the job-resource matchmaking
- see available resources via *condor\_status -constraint 'partitionableSlot' -af:t Name CPUs Memory Disk GPUs*
- match against a specific machine *condor\_q -better-analyse JOBID -reverse -machine MACHINE\_NAME*

# External Resources

- ETP has access to external computing resources via its batch system
  - Throughput Optimized Analysis System (TOpAS) at GridKa
    - use `condor_status -pool f03-001-140-e.gridka.de` to for TOpAS
  - BWForCluster NEMO2 at Uni. of Freiburg
    - create a NEMO2 account and add an acknowledgment in your thesis, see <https://wiki.etp.kit.edu/books/computing-and-infrastructure/page/nemo-cluster>
- external resources are only available with `+RemoteJob = True` in the JDL file
- `/home`, `/work`, or `/ceph` from ETP are not mounted on external resources  
⇒ transfer data via HTCondor (< 10 MB per job) or Grid protocols (XRootD, WebDAV, ...)

# Transfer Data via HTCondor

- transfers via htcondor runs through the scheduler ⇒ inefficient and extra load on portal machines

# Transfer Data via HTCondor

- transfers via htcondor runs through the scheduler ⇒ inefficient and extra load on portal machines
- transfer files to the WN: *transfer\_input\_files*
- transfer files from the WN at the end of the job: *transfer\_output\_files = output.root*
- rename output files: *transfer\_output\_remaps = "output.root = output/output\_\$(Process).root"*

```
Universe = container
container_image = /cvmfs/unpacked.cern.ch/registry.hub.docker.com/cverstege/alma9-gridjob:latest
executable = ./executable.sh

request_cpus = 1
request_memory = 2000
request_disk = 500000

transfer_input_files = /work/$(whoami)/HTCondor_Tutorial/01_simple_submit/mc_example.py

# DO NOT DO THAT FOR FILES BIGGER THAN 10MB
transfer_output_files = output.root
transfer_output_remaps = "output.root = output_$(Process).root"

log = log_$(Process).log
stdout = stdout_$(Process).log
stderr = stderr_$(Process).log

accounting_group=belle

queue
```

# Data intensive Jobs

- jobs with more than 2 GB input files per 30 min runtime are data intensive

# Data intensive Jobs

- jobs with more than 2 GB input files per 30 min runtime are data intensive
- write on local storage on the WN and copy the results at the end of the jobs

# Data intensive Jobs

- jobs with more than 2 GB input files per 30 min runtime are data intensive
- write on local storage on the WN and copy the results at the end of the jobs
- for some Grid storage requires a **VOMS (Virtual Organisation Membership Service) proxy**
  - location of the voms-proxy usually at `/tmp/x509up_u$(id -u)`
  - location can be set via environment variable `$X509_USER_PROXY`
  - HTCondor copies your proxy to the WN and set `$X509_USER_PROXY`

# Data intensive Jobs

- jobs with more than 2 GB input files per 30 min runtime are data intensive
- write on local storage on the WN and copy the results at the end of the jobs
- for some Grid storage requires a **VOMS (Virtual Organisation Membership Service) proxy**
  - location of the voms-proxy usually at `/tmp/x509up_u$(id -u)`
  - location can be set via environment variable `$X509_USER_PROXY`
  - HTCondor copies your proxy to the WN and set `$X509_USER_PROXY`

```
Universe = container
container_image = /cvmfs/unpacked.cern.ch/registry.hub.docker.com/cverstege/alma
executable = ./executable.sh
request_cpus = 1
request_memory = 2000
request_disk = 500000
transfer_input_files = /work/$(whoami)/HTCondor_Tutorial/01_simple_submit/mc_ex
# DO NOT DO THAT FOR FILES BIGGER THAN 10MB
transfer_output_files = output.root
transfer_output_remaps = "output.root=_output_$(ProcID).root"
log = log_(ProcID).log
stdout = stdout_(ProcID).log
stderr = stderr_(ProcID).log
accounting_group=belle
queue
```

# Transfer Data via XRootD

- files can be copied, streamed and written via XRootD
- `/ceph/srv`
  - only readable from ETP and TOpAS
  - via `root://ceph-node-j.etp.kit.edu://<user>`  
e.g. `xrdcp root://ceph-node-j.etp.kit.edu://mschnepf/testfile.txt file:///tmp/testfile.txt`

# Transfer Data via XRootD

- files can be copied, streamed and written via XRootD
- `/ceph/srv`
  - only readable from ETP and TOpAS
  - via `root://ceph-node-j.etp.kit.edu://<user>`  
e.g. `xrdcp root://ceph-node-j.etp.kit.edu://mschnepf/testfile.txt file:///tmp/testfile.txt`
- NRG at GridKa
  - CMS `root://cmsdcache-kit-disk.gridka.de:1094//store/user/<user>`
  - Belle `root://dcache-kit.gridka.de:1094//pnfs/gridka.de/belle/disk-only/LOCAL/user/<user>`
  - voms proxy required

# Transfer Data via XRootD

- files can be copied, streamed and written via XRootD
- `/ceph/srv`
  - only readable from ETP and TOpAS
  - via `root://ceph-node-j.etp.kit.edu://<user>`  
e.g. `xrdcp root://ceph-node-j.etp.kit.edu://mschnepf/testfile.txt file:///tmp/testfile.txt`
- NRG at GridKa
  - CMS `root://cmsdcache-kit-disk.gridka.de:1094//store/user/<user>`
  - Belle `root://dcache-kit.gridka.de:1094//pnfs/gridka.de/belle/disk-only/LOCAL/user/<user>`
  - voms proxy required
- TOpAS and other worker nodes can cache files via XRootD
  - enables fast processing
  - files are identified by filename e.g. `/mschnepf/testfile.txt`  
⇒ do not overwrite files, the cache does not update

# Transfer Data via XRootD

- files can be copied, streamed and written via XRootD
- `/ceph/srv`
  - only readable from ETP and TOpAS
  - via `root://ceph-node-j.etp.kit.edu://<user>`  
e.g. `xrdcp root://ceph-node-j.etp.kit.edu://mschnepf/testfile.txt file:///tmp/testfile.txt`
- NRG at GridKa
  - CMS `root://cmsdcache-kit-disk.gridka.de:1094//store/user/<user>`
  - Belle `root://dcache-kit.gridka.de:1094//pnfs/gridka.de/belle/disk-only/LOCAL/user/<user>`
  - voms proxy required
- TOpAS and other worker nodes can cache files via XRootD
  - enables fast processing
  - files are identified by filename e.g. `/mschnepf/testfile.txt`  
⇒ do not overwrite files, the cache does not update
- most root installations can open files from remote
  - e.g. `myroot.py root://ceph-node-j.etp.kit.edu://mschnepf/test_rootfile.txt`

# Send Software Sandbox via Grid Storage

- code can run outside of ETP
  - no absolute paths
  - `/work`, `/ceph`, `/home` is not available
- pack software into one tarball-file (sandbox)
- copy sandbox to Grid storage
- job downloads and unpacks sandbox
- some software needs to linked at new location, e.g. CMSSW

```
##### WN information
echo -e "Hostname_\$(hostname)"
echo -e "user_\$(whoami)"
echo -e "spawndir_\$(pwd)"

##### setup environment
## download sandbox
xrdcp root://dcache-kit.gridka.de:1094//pnfs/gridka.de/belle/disk-only/LHCb/CMSSW_15_0_10_patch3.tar.gz /tmp/
tar -xf /tmp/sandbox_test.tar
## setup sandbox software
source /cvmfs/cms.cern.ch/cmsset_default.sh
cd CMSSW_15_0_10_patch3/src
scram b ProjectRename
eval `scramv1 runtime -sh`

##### start workload at spawndir and create output.root with results
cd ${spawndir}
which combine

##### copy output from WN
#gfal-copy output.root
root://dcache-kit.gridka.de:1094://pnfs/gridka.de/belle/disk-only/LHCb/CMSSW_15_0_10_patch3/output.root
```

# GPUs

- only TOpAS provides GPUs  $\Rightarrow +RemoteJob=True$
- use the docker image (*container\_image = docker://...*)
- request a GPU via *RequestGPU=1* in the JDL
- check if your job can run at TOpAS: *condor\_q -better-analyse -pool f03-001-140-e.gridka.de JOBID*
- CUDA is not provided by the system

# GPUs

- only TOpAS provides GPUs ⇒ *+RemoteJob=True*
- use the docker image (*container\_image = docker://...*)
- request a GPU via *RequestGPU=1* in the JDL
- check if your job can run at TOpAS: *condor\_q -better-analyse -pool f03-001-140-e.gridka.de JOBID*
- CUDA is not provided by the system
- setup environment
  - source an existing environment, e.g., CMSSW / BASF2 / LCG-Stack
  - sandbox via Grid storage
  - setup your environment via pip or conda
  - build your own container and publish it on docker hub and CVMFS-unpacked

```
#!/bin/bash
echo "hostname_\$(hostname)"

# Source LCG-Stack (some of them do not work)
source /cvmfs/sft.cern.ch/lcg/views/LCG_107_cuda/x86_64-el9-gcc11
-opt/setup.sh

env
python3 mnist_training.py

ls -lh
```

# GPUs

- only TOpAS provides GPUs  $\Rightarrow +RemoteJob=True$
- use the docker image (*container\_image = docker://...*)
- request a GPU via *RequestGPU=1* in the JDL
- check if your job can run at TOpAS: *condor\_q -better-analyse -pool f03-001-140-e.gridka.de JOBID*
- CUDA is not provided by the system
- setup environment
  - source an existing environment, e.g., CMSSW / BASF2 / LCG-Stack
  - sandbox via Grid storage
  - setup your environment via pip or conda
  - build your own container and publish it on docker hub and CVMFS-unpacked
- request parts of a GPU via *+RequestGPUMemoryMB = 2000*
  - GPU processing power is still shared
  - works only with CUDA software
- **DO NOT SET *NVIDIA\_VISIBLE\_DEVICES***; HTCondor does it

```
#!/bin/bash
echo "hostname_\$(hostname)"

# Source LCG-Stack (some of them do not work)
source /cvmfs/sft.cern.ch/lcg/views/LCG_107_cuda/x86_64-el9-gcc11
-opt/setup.sh

env
python3 mnist_training.py
ls -lh
```

# Further Information

- ETP Wiki: HTCondor at ETP
- ETP Batch System monitoring
- HTCondor youtube user tutorials
- ReadTheDocs
- ETP mattermost: etp-htcondor
- ETP mattermost: etp-htcondor-gpus

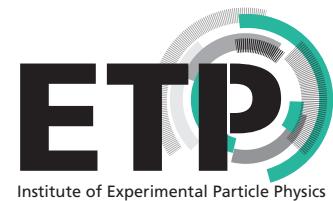




# WE WANT YOU!

<https://www.clker.com/clipart-753080.html>

- batch admin
- computing work
  - Bachelorthesis
  - Masterthesis
  - PhD Topics
  - Hiwi



# Backup

# CMS Example Analysis Workflow (2017 data)

